

Project 2 MLSec

—

Universal adversarial attack against three models

- **Names:** Marcel Plocher, Moritz Aberle
 - **Matricola:** IA/65928, IA/65830
 - **Instructor:** Battista Biggio

Outline

- Problem and Requirements
- Choosing Models from RobustBench
- Theoretical Approach
- Combining the Models
- Generation of Adversarial Examples
- Results and Transferability
- Challenges and Conclusion

Problem and Requirements

1. Generate universal adversarial examples
 - Choose 3 models from RobustBench
 - Using the CIFAR10 dataset and the L-inf-norm
 - Untargeted
 - Universal – for all three models at the same time
2. Evaluated transferability
 - 7 other models
 - Using the same adversarial examples

Choosing Models from RobustBench

- Aspects
 - Less adversarial training
 - Non robust pretraining (e.g. Reinforcement Learning)
 - Are trained with weak or outdated robustness techniques
- 10 Models
 - Hendrycks2019Using
 - Chen2020Adversarial
 - Wong2020Fast
 - Engstrom2019Robustness
 - Ding2020MMA
 - Rice2020Overfitting
 - Huang2020Self
 - Sehwal2021Proxy_R18
 - Rebuffi2021Fixing_70_16_cutmix_extra
 - 3. Rade2021Helper_extra

Theoretical Approaches – PGD Attack

- Constraint: Maximum-confidence

$$\min_{\|\delta\| \leq \epsilon} L(x + \delta, y; \theta)$$

- Projected Gradient Descent (PGD)

$$x^* = x + \epsilon \operatorname{sign}(\nabla L(x, y, \theta))$$

$$\max_{\|\delta\|_\infty} L(x + \delta, y, \theta)$$

$$\delta^* = \epsilon \operatorname{sign}(\nabla L)$$

Combining the models

- **Method1:** Averaging of model outputs

- Computationally expensive
- Stable and deterministic

$$\frac{1}{M} \sum_{i=1}^M f_i(x)$$

- **Method2:** Taking the most frequent prediction (Not tracked by us)

- Simple and efficient method
- Deterministic

- **Method3:** Random model selection (uniform distribution)

- The attack generalizes better
- Computationally efficient
- Non-deterministic


$$F(x) = f_j(x), \text{ where } j \sim \text{Uniform}(\{1, 2, \dots, M\})$$

Generation of Adversarial Examples

- Using the PGD attack from SecML
- Using DLR-loss (Difference of Logits Ratio) instead of cross-entropy-loss
- Hyperparameter
 1. **Epsilon = 8/255**: Default value from RobustBench
 2. **N = 100**: Number of PGD-iterations
 3. **Alpha = Eps/N or 1/255**: Amount of perturbation at each step (how aggressive)
 4. **L-inf Norm**: Limits the maximum change applied to any pixel

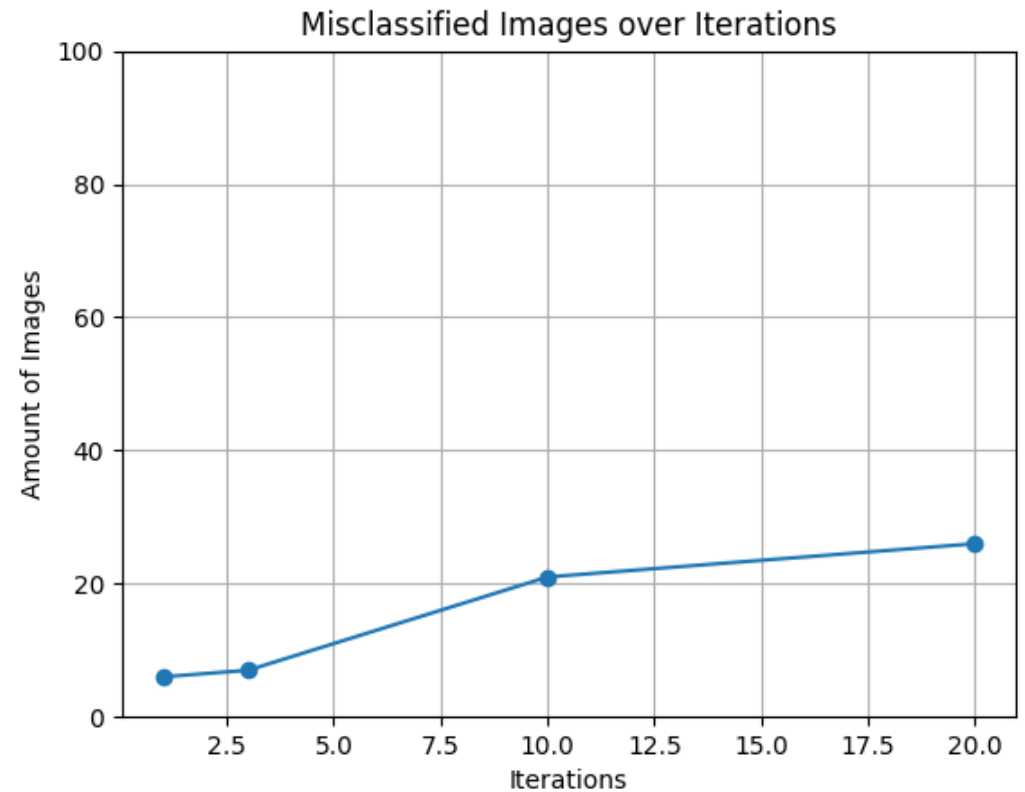
Results – Misclassified Images (Epsilon/N)

- Amount of misclassified images over all 3 models (epsilon/N)
 - 1 Iteration = 18/100 Images
 - 3 Iterations = 20/100 Images
 - 10 Iterations = 21/100 Images
 - 20 Iterations = 22/100 Image





Results – Misclassified Images (1/255)



- Amount of misclassified images over all 3 models (epsilon/N)
 - 1 Iteration = 6/100 Images
 - 3 Iterations = 7/100 Images
 - 10 Iterations = 21/100 Images
 - 100 Iterations = 26/100 Images





Results - Transferability ($N=1$, $\text{Alpha}=\text{Eps}/N$)




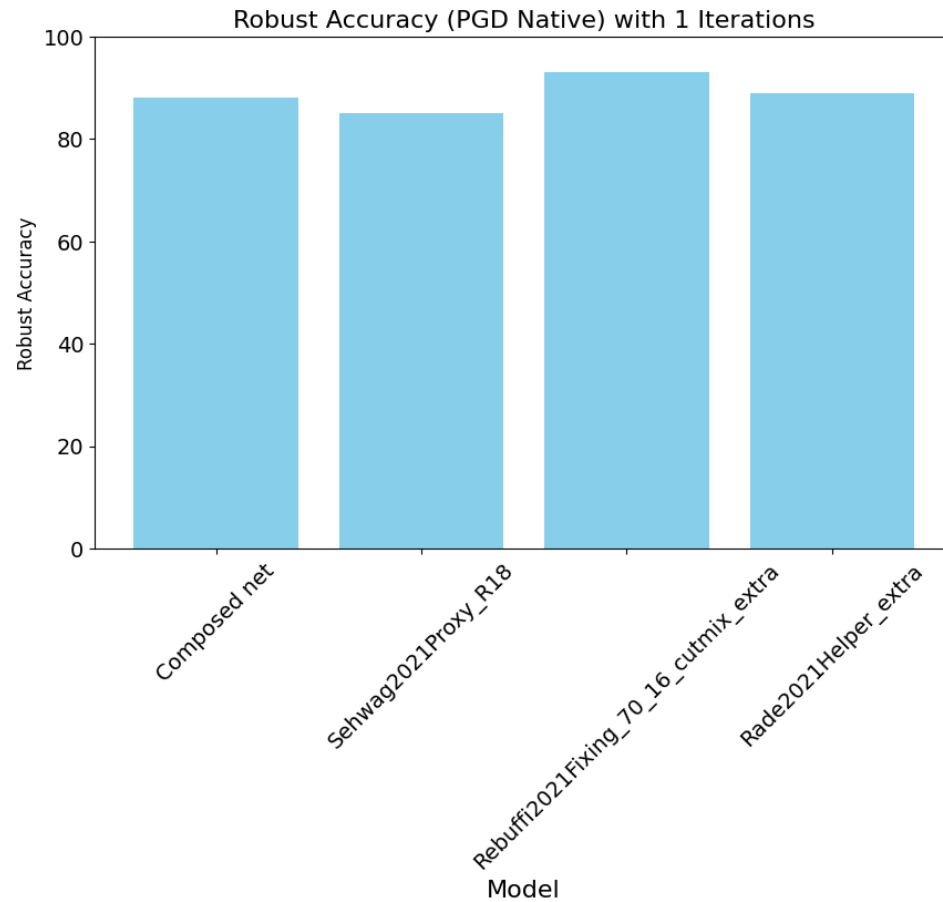
Results – Transferability (N=3, Alpha=Eps/N)





Results - Transferability (N=10, Alpha=Eps/N)





Results - Transferability (N=1, Alpha=1/255)





Results - Transferability (N=3, Alpha=1/255)



Results - Transferability (N=10, Alpha=1/255)



Results - Transferability (N=100, Alph=1/255)



Challenges and Conclusion

- With combination of 3 models the attack generalizes very good
- Finding the optimal hyperparameters is hard
 - Long runtime of the attack
 - Depending on concrete conditions (i.e. choosing epsilon)
 - Finding the optimal loss-function
 - Understanding the connection between the step size and the number of iterations